

High Performance Computing

Sony Dodla

2303A51666

Batch-23

Assignment 3: Load Balancing with Irregular Workloads

Scenario

An image processing pipeline processes images with different resolutions, leading to uneven computation time.

Objective

Study load imbalance in parallel loops.

Tasks

- 1.Simulate variable workload per iteration.
- 2.Parallelize using prange.
- 3.Observe execution time variation.
- 4.Discuss limitations of scheduling in Python.

Learning Outcomes

- Load imbalance
- Comparison with OpenMP scheduling
- Practical HPC limitations in Python

Python code

```
import numpy as np
import time
from numba import njit, prange
np.random.seed(0)
n_images = 50
# Simulate different image resolutions (irregular workload)
resolutions = np.random.randint(50_000, 500_000, size=n_images)
# -----
# Simulate image processing workload
```

```
# -----
@njit
def process_image(resolution):
    """
    Artificially heavy computation proportional to image resolution.
    """

    acc = 0.0

    for i in range(resolution):
        acc += np.sqrt(i * 0.5)

    return acc

# -----
# Serial version
# -----
@njit
def serial_processing(resolutions):
    results = np.zeros(len(resolutions))

    for i in range(len(resolutions)):
        results[i] = process_image(resolutions[i])

    return results

# Serial timing
t0 = time.time()
serial_processing(resolutions)
t1 = time.time()

print("Number of images:", n_images)
print("Min resolution:", resolutions.min())
print("Max resolution:", resolutions.max())
print(f"Serial execution time : {t1 - t0:.4f} s")

# -----
# Parallel version using prange
# -----
@njit(parallel=True)
```

```

def parallel_processing(resolutions):
    results = np.zeros(len(resolutions))
    for i in prange(len(resolutions)):
        results[i] = process_image(resolutions[i])
    return results

# Parallel timing
t2 = time.time()
parallel_processing(resolutions)
t3 = time.time()
print(f"Parallel execution time : {t3 - t2:.4f} s")

```

Output:

Number of images: 50
 Min resolution: 57877
 Max resolution: 498242
 Serial execution time : 0.2397 s
 Parallel execution time : 0.3970 s

C code:

```

#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#include <omp.h>
#define N_IMAGES 1000
#define MAX_THREADS 64
double process_image(int resolution)
{
    double acc = 0.0;
    for (int k = 0; k < resolution; k++)
        acc += sqrt(k * 0.5);
    return acc;

```

```

}

void run_schedule(int *resolution, omp_sched_t sched_type, const char *name)
{
    int i;
    double total = 0.0;
    int nthreads = omp_get_max_threads();
    int work[MAX_THREADS] = {0};
    double t_start[MAX_THREADS] = {0.0};
    double t_end[MAX_THREADS] = {0.0};
    omp_set_schedule(sched_type, 0);
    double start = omp_get_wtime();
    #pragma omp parallel private(i)
    {
        int tid = omp_get_thread_num();
        t_start[tid] = omp_get_wtime();
        #pragma omp for schedule(runtime) reduction(+:total)
        for (i = 0; i < N_IMAGES; i++) {
            process_image(resolution[i]);
            work[tid]++;
        }
        t_end[tid] = omp_get_wtime();
    }
    double end = omp_get_wtime();
    printf("\n== %s scheduling ==\n", name);
    printf("Total execution time: %f s\n", end - start);
    printf("Per-thread workload:\n");
    for (i = 0; i < nthreads; i++) {
        printf("Thread %2d | images: %4d | active time: %6.3f s\n",
               i, work[i], t_end[i] - t_start[i]);
    }
}

```

```

int main()
{
    int resolution[N_IMAGES];
    srand(0);
    for (int i = 0; i < N_IMAGES; i++)
        resolution[i] = 10000 + rand() % 200000;
    printf("Images : %d\n", N_IMAGES);
    printf("Threads : %d\n", omp_get_max_threads());
    run_schedule(resolution, omp_sched_static, "STATIC");
    run_schedule(resolution, omp_sched_dynamic, "DYNAMIC");
    run_schedule(resolution, omp_sched_guided, "GUIDED");
    return 0;
}

```

Output:

Images : 1000

Threads : 4

==== STATIC scheduling ===

Total execution time: 0.466385 s

Per-thread workload:

Thread 0 | images: 250 | active time: 0.454 s

Thread 1 | images: 250 | active time: 0.441 s

Thread 2 | images: 250 | active time: 0.451 s

Thread 3 | images: 250 | active time: 0.457 s

==== DYNAMIC scheduling ===

Total execution time: 0.427426 s

Per-thread workload:

Thread 0 | images: 191 | active time: 0.427 s

Thread 1 | images: 177 | active time: 0.423 s

Thread 2 | images: 320 | active time: 0.420 s

Thread 3 | images: 312 | active time: 0.427 s

==== GUIDED scheduling ===

Total execution time: 0.497476 s

Per-thread workload:

Thread 0 | images: 233 | active time: 0.497 s

Thread 1 | images: 305 | active time: 0.496 s

Thread 2 | images: 185 | active time: 0.496 s

Thread 3 | images: 277 | active time: 0.438 s